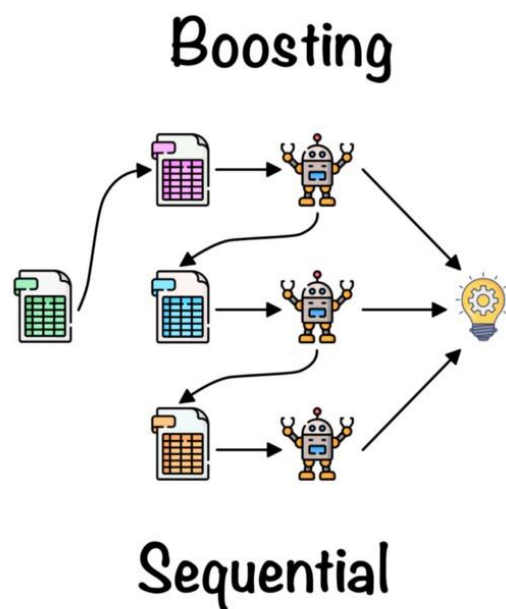



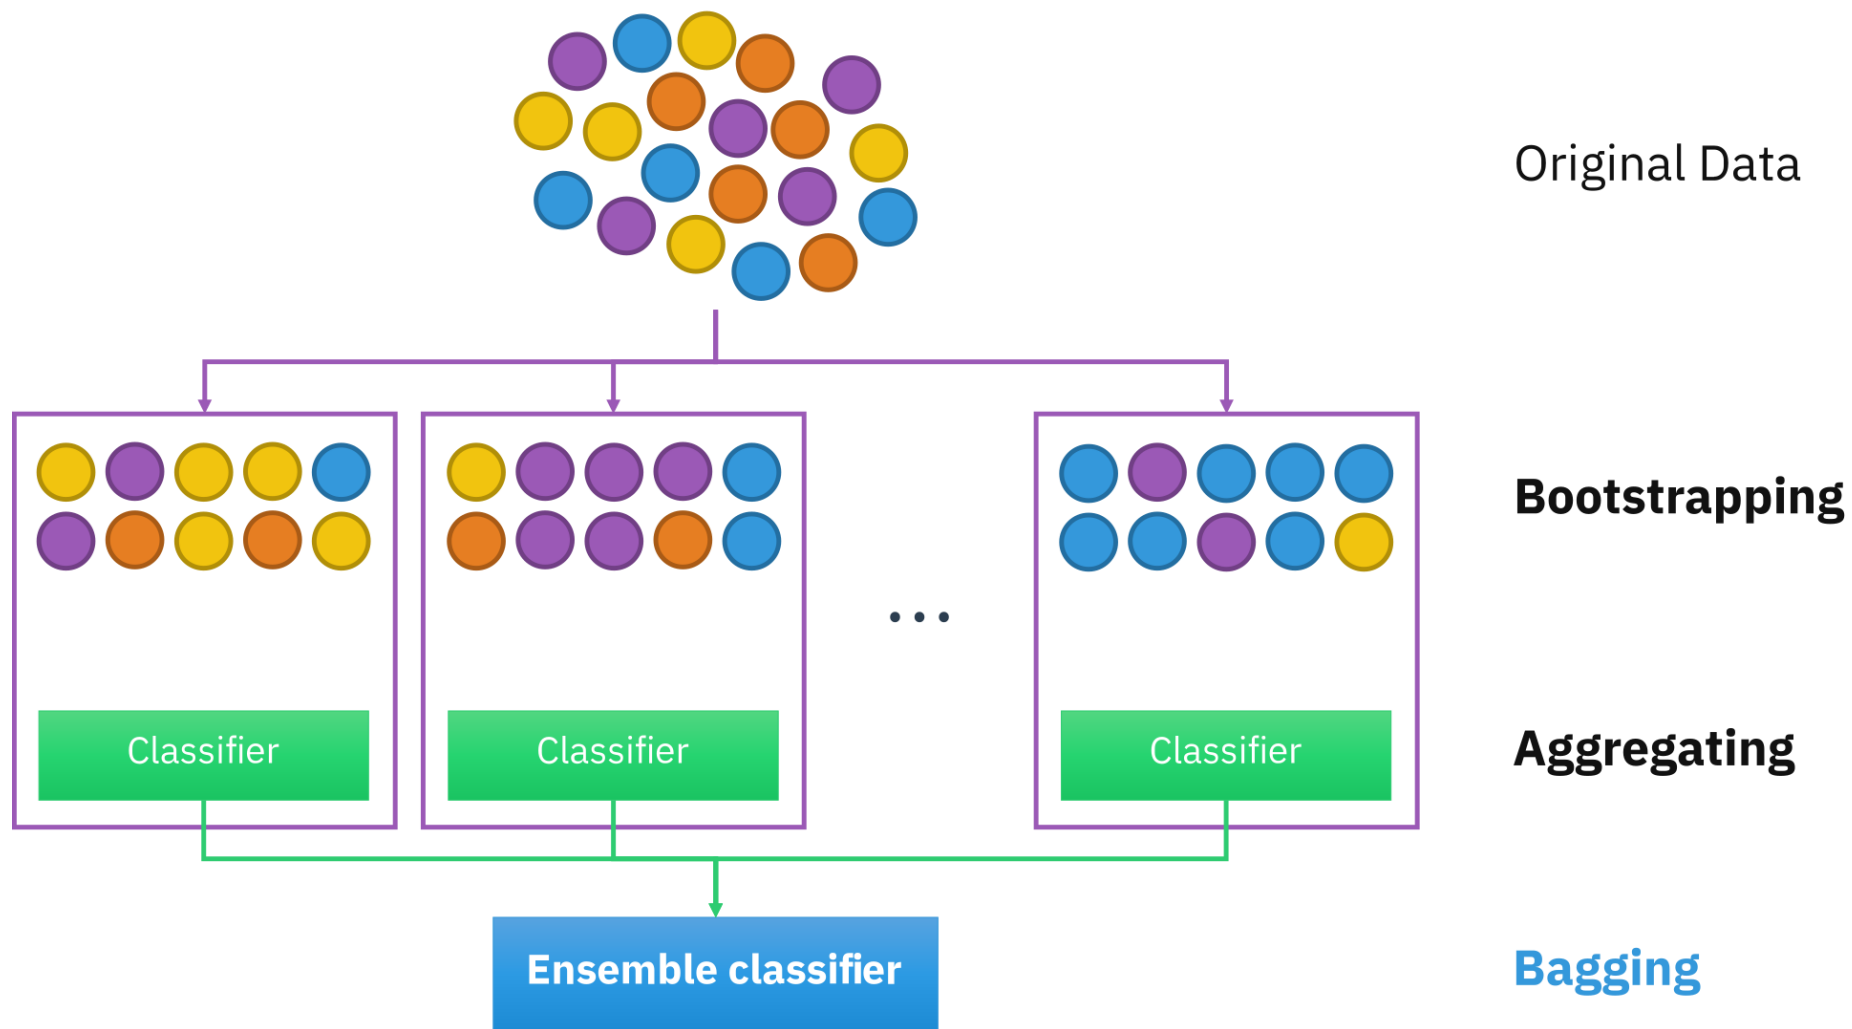
Bagging and Boosting



赵海涛

haitaozhao@ecust.edu.cn


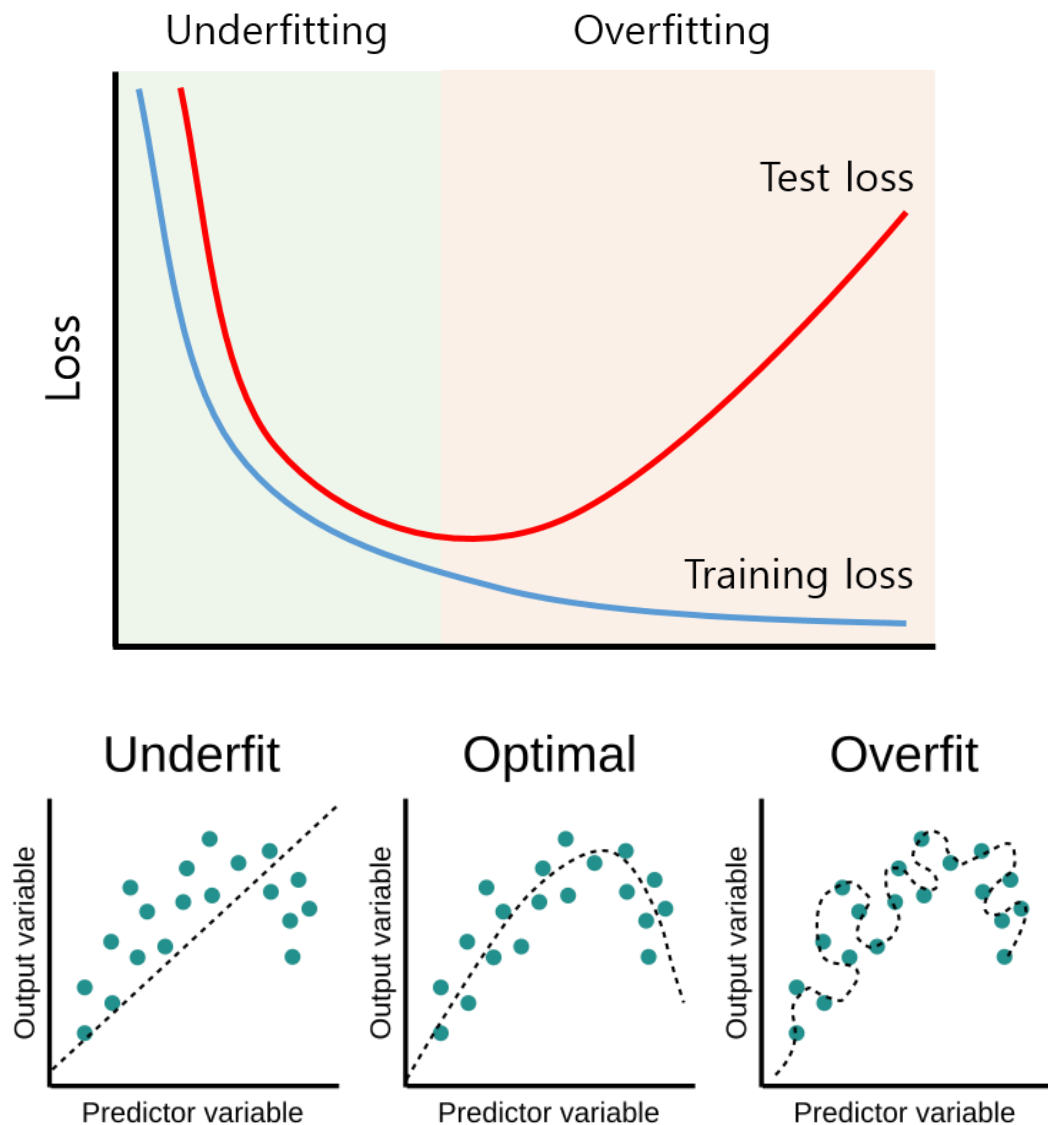
开课一张图



大纲

- 偏差与方差
- Bagging
- Boosting

回归的偏差和方差



回归的偏差和方差

- 假定 $\phi(x)$ 是一个未知函数，希望通过 $\phi(x)$ 产生的 N 个样本的数据集 \mathcal{D} 来估计 $\phi(x)$. 假设候选函数属于集合 \mathcal{F} . 令 $f(x; \mathcal{D})$ 是被估计函数，则 $f(x; \mathcal{D}) \in \mathcal{F}$.
- 针对固定的 N ，对数据集 \mathcal{D} 求平均

$$E_{\mathcal{D}}[(f(x; \mathcal{D}) - \phi(x))^2] = \underbrace{(E_{\mathcal{D}}[f(x; \mathcal{D})] - \phi(x))^2}_{\text{bias}^2} + \underbrace{E_{\mathcal{D}}[(f(x; \mathcal{D}) - E_{\mathcal{D}}[f(x; \mathcal{D})])^2]}_{\text{variance}}.$$

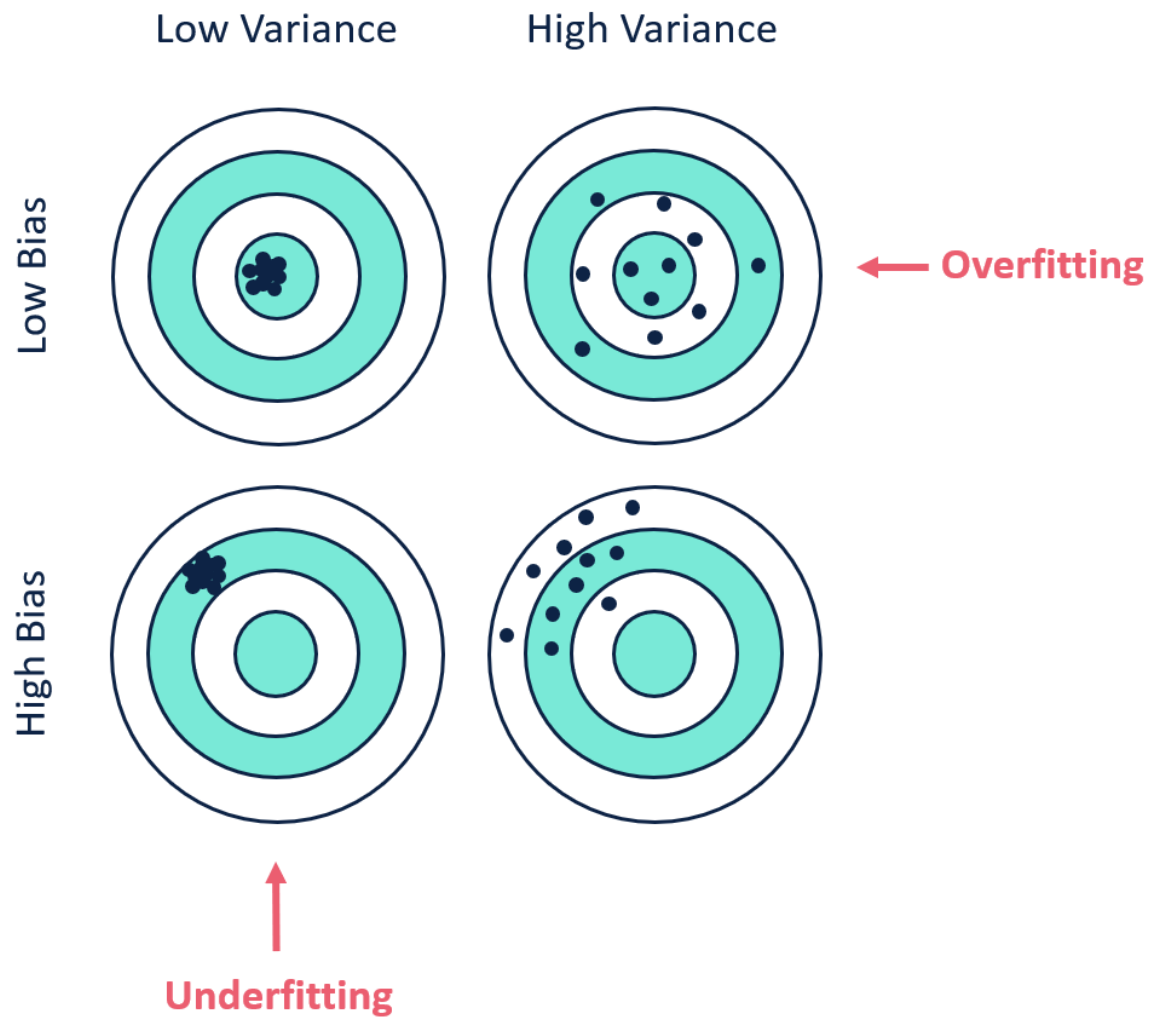
- 令 $y = \phi(x) + \varepsilon$, 其中 $E(\varepsilon) = 0, \text{var}(\varepsilon) = \sigma_{\varepsilon}^2$. $\phi(x)$ 和噪声 ε 独立， 则

$$\begin{aligned} \text{Error}(x) &\triangleq E_{\mathcal{D}}[(y - f(x; \mathcal{D}))^2] \\ &= \sigma_{\varepsilon}^2 + \text{bias}^2 + \text{variance}. \end{aligned}$$


回归的偏差和方差

- 大体来说，被估计函数的泛化误差是其在测试集上的误差期望
- 对一个固定的回归算法，针对一个新采样的训练集，被估计的函数 f 一般不相同。方差 (*variance*) 表示 f 对训练集改变的敏感性
- 如果 \mathcal{F} 中的函数都过于简单，则候选函数不能很好的拟合数据，产生偏差 (*bias*)。偏差可以是：
 - 确定性的 (hard): 没有 $f \in \mathcal{F}$ 能够拟合数据；
 - 随机的 (soft): $f \in \mathcal{F}$ 能够拟合数据的先验概率非常小


回归的偏差和方差



回归的偏差和方差



the bias vs. variance trade-off



如何估计方差

$$E_{\mathcal{D}}[(f(x; \mathcal{D}) - E_{\mathcal{D}}[f(x; \mathcal{D})])^2]$$

variance

- 数据独立同分布
- 如果可以得到大量不同的训练集将是非常好的。这种情况下，可以针对每个训练集进行估计，并对被估计函数的分布进行推断
- 通常额外的训练集并不容易获得

Bootstrap

- 通过从原始样本中重新抽样来估计估计量的抽样分布。[Efron, The Annals of Statistics, 1979]
- 抽样是根据经验分布进行的

	Samples	Estimator
$P \Rightarrow x_1, x_2, \dots, x_n \xrightarrow{\hat{p}}$	$z_{11}, z_{12}, z_{13}, \dots, z_{1n}$	e_1
	$z_{21}, z_{22}, z_{23}, \dots, z_{2n}$	e_2
	\vdots	\vdots
	$z_{m1}, z_{m2}, z_{m3}, \dots, z_{mn}$	e_m

主要思想

- $(x_1, x_2, \dots, x_n) \sim P$. 注意: 分布函数 P 是未知的
- 采样 m 个数据集 Y_1, Y_2, \dots, Y_m . $Y_i = (z_{i1}, z_{i2}, \dots, z_{in})$ 包含 n 个从经验分布中抽取的样本, 训练集的经验分布为:

$$\Pr[z_{jk} = x_i] = \frac{\#x_i}{n}$$

其中 $\#x_i$ 表示 x_i 在原始训练集中出现的次数

- $Y_i \sim \hat{P}$
- $\hat{P} = P$ 在理论上并不能保证。但 \hat{P} 是 P 的一个近似。

Bagging

- Bagging 可以减少方差
- 假设我们要估计的是某个参数，参数为 θ ，并假定 $f_{ag} = E_P \hat{f}^*(x)$.
估计函数 $\hat{f}^*(x)$ 是基于样本集 $Y^* = (z_1^*, z_2^*, \dots, z_n^*) \sim P$ ，则

$$\begin{aligned} E_P[\theta - \hat{f}^*(x)]^2 &= E_P[\theta - f_{ag} + f_{ag} - \hat{f}^*(x)]^2 \\ &= E_P[\theta - f_{ag}]^2 + E_P[f_{ag} - \hat{f}^*(x)]^2 \\ &\geq E_P[\theta - f_{ag}]^2 \end{aligned}$$

Bagging Averages (Bootstrap aggregation)

- 假设我们有独立的训练集 $\mathcal{D}_1, \mathcal{D}_2, \dots, \mathcal{D}_M$. g_i 是在 \mathcal{D}_i ($i = 1, 2, \dots, M$) 上得到的分类器, 令

$$f = E[g] = \frac{1}{M} \sum_{k=1}^M g_i.$$

- 平均可以得到更小的方差
- Bagging的步骤 (考虑二分类):
 - 针对训练集 \mathcal{D} , 独立采样 $\mathcal{D}_1, \dots, \mathcal{D}_M$
 - 在 \mathcal{D}_i ($i = 1, 2, \dots, M$)上分别估计 g_i
 - 得到最终的分类器

$$f(x) = \text{sign} \frac{1}{M} \sum_{k=1}^M g_i(x)$$

Example: 随机森林

Algorithm 15.1 *Random Forest for Regression or Classification.*

1. For $b = 1$ to B :
 - (a) Draw a bootstrap sample \mathbf{Z}^* of size N from the training data.
 - (b) Grow a random-forest tree T_b to the bootstrapped data, by recursively repeating the following steps for each terminal node of the tree, until the minimum node size n_{min} is reached.
 - i. Select m variables at random from the p variables.
 - ii. Pick the best variable/split-point among the m .
 - iii. Split the node into two daughter nodes.
2. Output the ensemble of trees $\{T_b\}_1^B$.

To make a prediction at a new point x :

Regression: $\hat{f}_{\text{rf}}^B(x) = \frac{1}{B} \sum_{b=1}^B T_b(x)$.

Classification: Let $\hat{C}_b(x)$ be the class prediction of the b th random-forest tree. Then $\hat{C}_{\text{rf}}^B(x) = \text{majority vote } \{\hat{C}_b(x)\}_1^B$.

- The Elements of Statistical Learning. T. Hastie, R. Tibshirani and J. H. Friedman.

随机森林 (MATLAB Code)

```
%% bagged decision tree
```

```
load iris; Data = X;
```

```
[Projection, Projection2, rs] = PCA_f(3, 50, Data, ...  
Data);
```

```
trn.X = Projection(1:2,51:150);
```

```
trn.y =[ones(1,50) 2*ones(1,50)];
```

```
% random forest
```

```
b = TreeBagger(50,trn.X',trn.y,'OOBPred','on')
```

```
plot(oobError(b))
```

```
% evaluate the tree classifier
```

```
b = TreeBagger(15,trn.X',trn.y,'OOBPred','on')
```


```
[sfit,scores] = predict(b,trn.X'); % Find ...
```

```
assigned class numbers
```

```
y_pred = str2num(cell2mat(sfit));
```

```
crosstab(y(51:150),y_pred)
```

随机森林 (MATLAB Plot)



怎样组合弱分类器

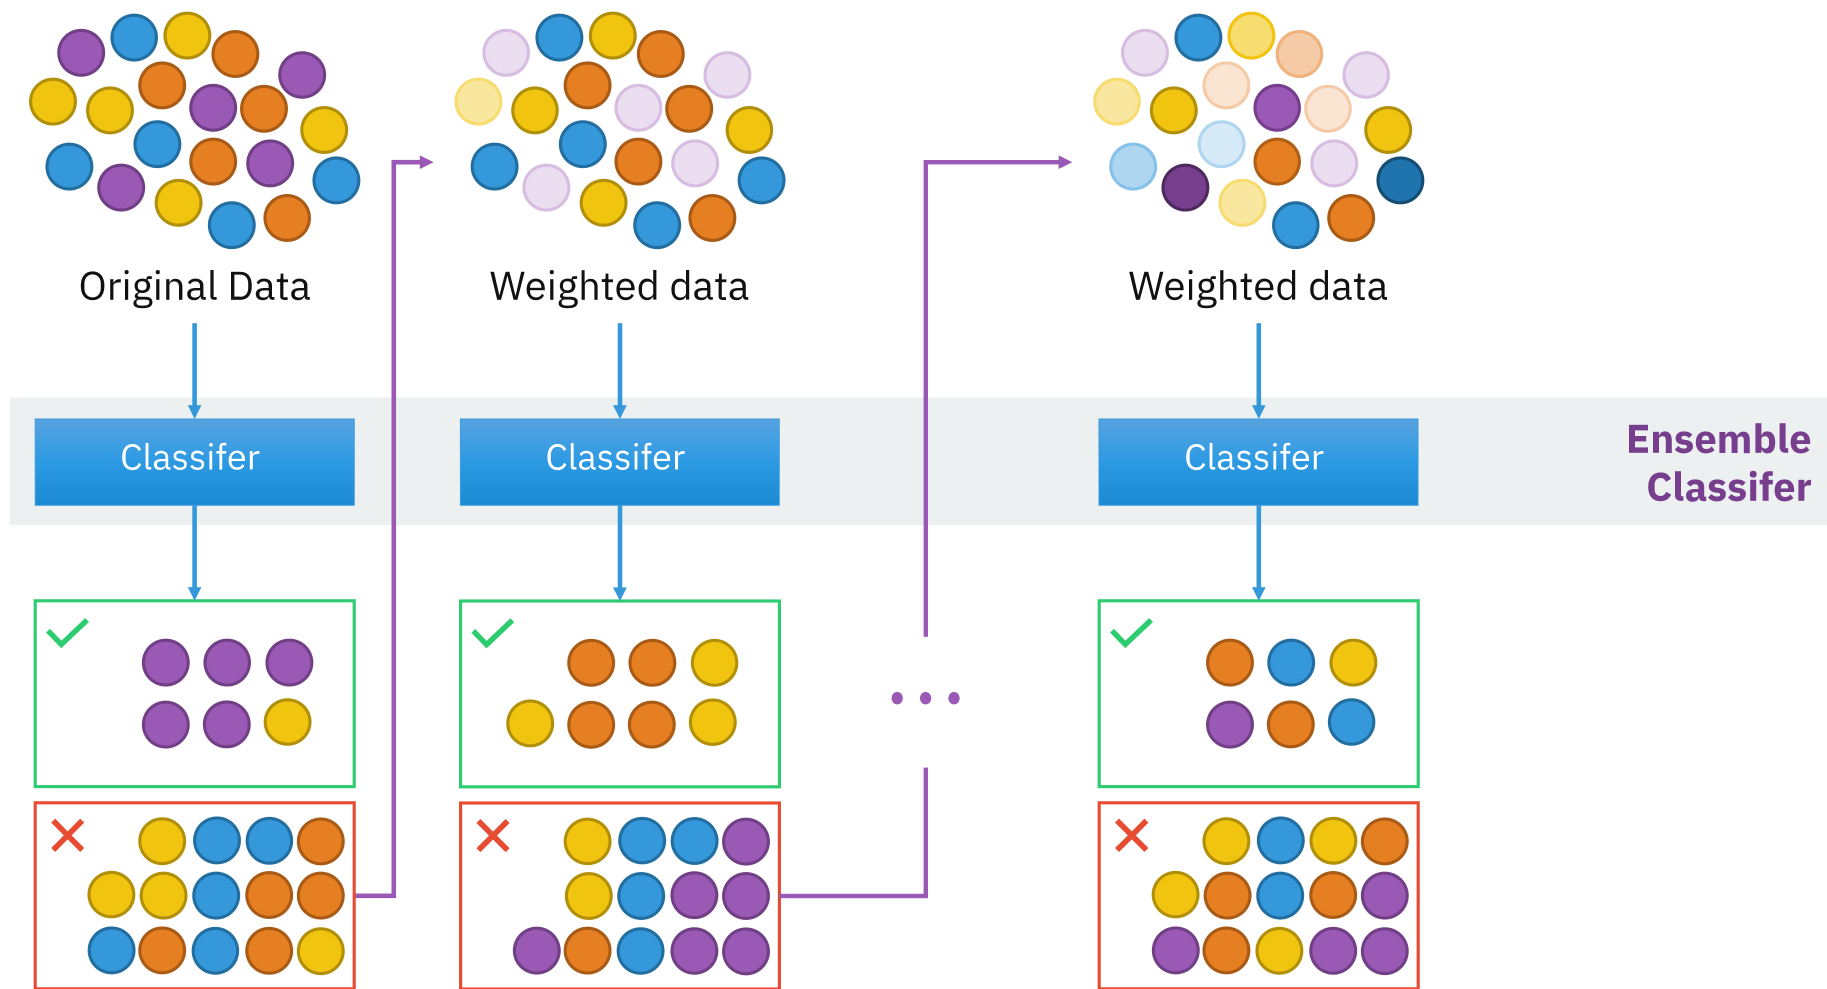
- 多专家组合

一种**并行**结构，**所有**的弱分类器都给出各自的预测结果，通过“组合器”把这些预测结果转换为最终结果。 eg. **投票（voting）** 及其变种、混合专家模型

- 多级组合

一种**串行**结构，其中下一个分类器只在前一个分类器预测不够准（不够自信）的实例上进行训练或检测。 eg. **级联算法（cascading）**

另一张图



提升 (Boosting)

$$(E_{\mathcal{D}}[f(x; \mathcal{D})] - \phi(x))^2$$

*bias*²

- Boosting 可以减小偏差.
- 基分类器族 \mathcal{G} 有很大的偏差（例如线性分类器、决策树桩），但学习总是产生分类器 g 比随机猜测更好（在训练集上）。
- Boosting 的先决条件:
 - 对于 $g \in \mathcal{G}$, g 的训练错误率有上界, 且 $0 < \widehat{\text{Error}}(g) \leq \delta < \frac{1}{2}$
 - 可以在加权的数据上进行训练

Boosting先決条件

$$\left. \begin{array}{l} h_1(x) \in \{-1, +1\} \\ h_2(x) \in \{-1, +1\} \\ \vdots \\ h_T(x) \in \{-1, +1\} \end{array} \right\}$$

Weak classifiers

slightly better than random

$$H_T(x) = \text{sign} \left(\sum_{t=1}^T \alpha_t h_t(x) \right)$$

strong classifier

Adaboost基本概念

两个问题如何解决：

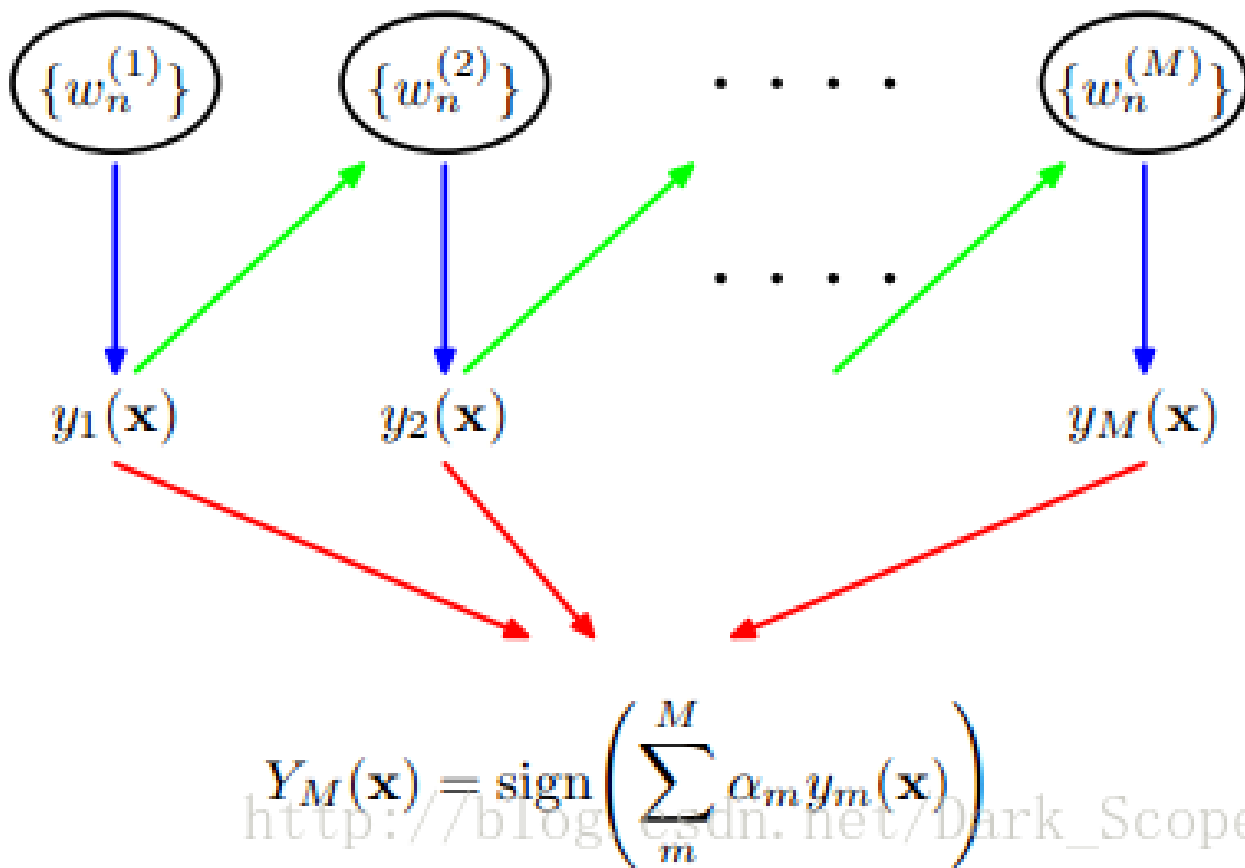
每一轮如何改变训练数据的权值或概率分布？

AdaBoost：提高那些被前一轮弱分类器错误分类样本的权值，降低那些被正确分类样本的权值

如何将弱分类器组合成一个强分类器？

AdaBoost：加权多数表决，加大分类误差率小的弱分类器的权值，使其在表决中起较大的作用，减小分类误差率大的弱分类器的权值，使其在表决中起较小的作用。

Adaboost基本概念



Adaboost.M1 (Freund and Schapire)

- Initialize the observation weights $w_i = \frac{1}{N}$, $i = 1, 2, \dots, N$
- For $m = 1$ to M :
- Fit a classifier $g_m(x)$ to the training data using weights w_i .
- Compute

$$\text{err}_m = \sum_{i=1}^N w_i I(y_i \neq g_m(x_i))$$

- Compute


$$\alpha_m = \frac{1}{2} \ln \left[\frac{1 - \text{err}_m}{\text{err}_m} \right]$$

- Set $w_i \leftarrow \frac{w_i \exp[-\alpha_m y_i g_m(x_i)]}{\sum_{i=1}^N w_i \exp[-\alpha_m y_i g_m(x_i)]}$
- Output

$$H_g(x) = \text{sign} \left[\sum_{m=1}^M \alpha_m g_m(x) \right]$$


Adaboost图解

用决策桩（就是阈值判别器）去划分平面上红色、蓝色小球，初始状态是这样的：




Adaboost

然后第一个桩分类器竖着划分，错分了两个：




于是在重新计算分布后，它们的权重变大了：

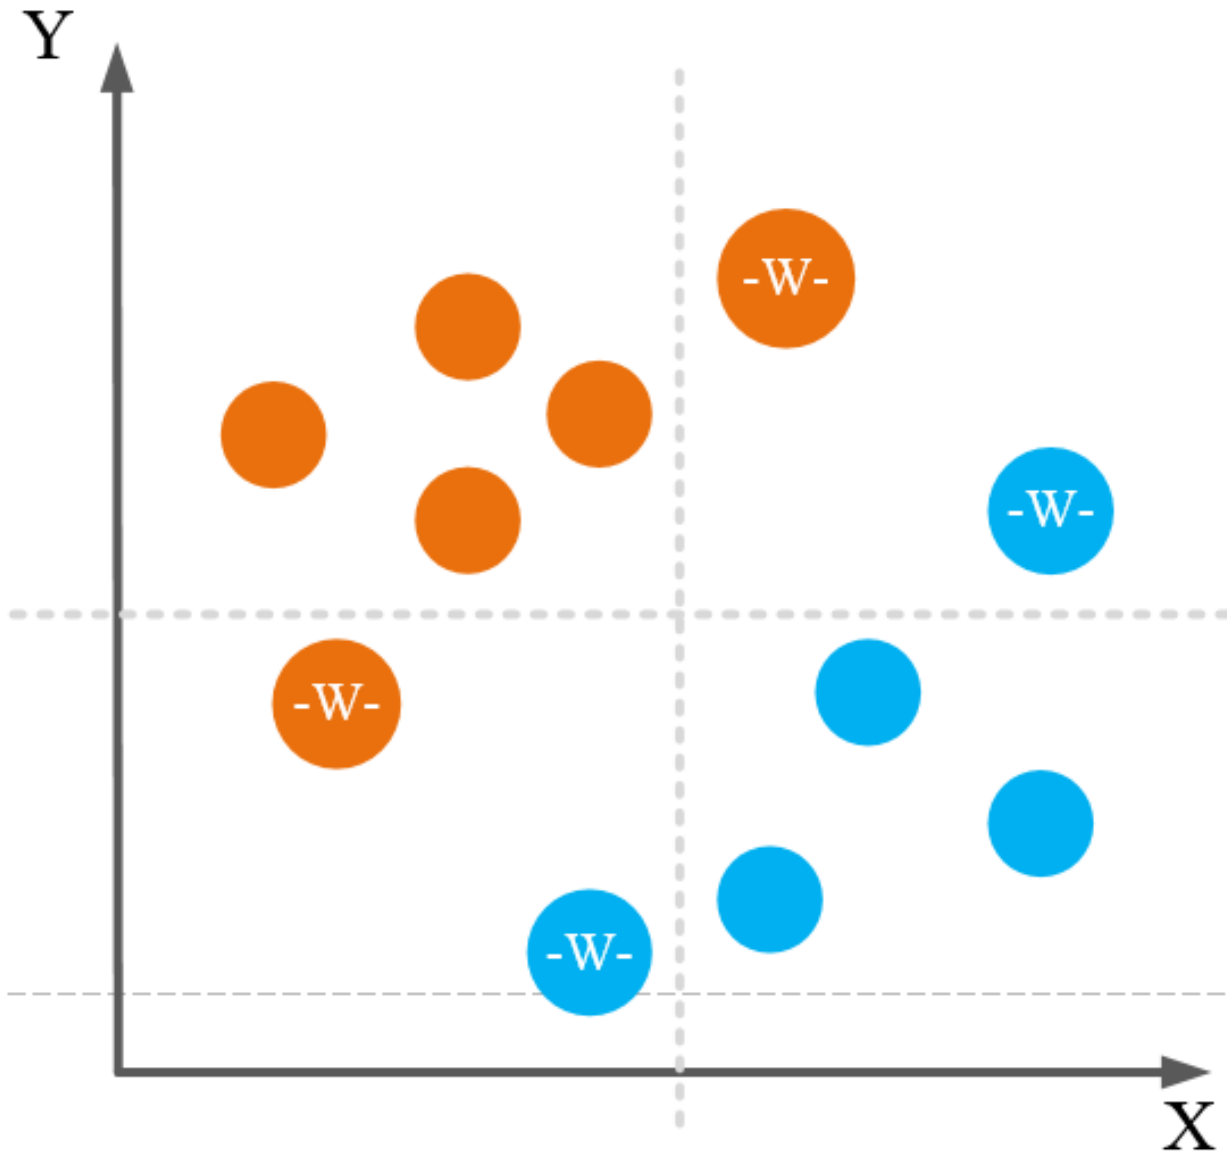


Adaboost

第二个分类器更多的考虑了被错分的样本，
然而第一次被分对的样本又被分错了两个：

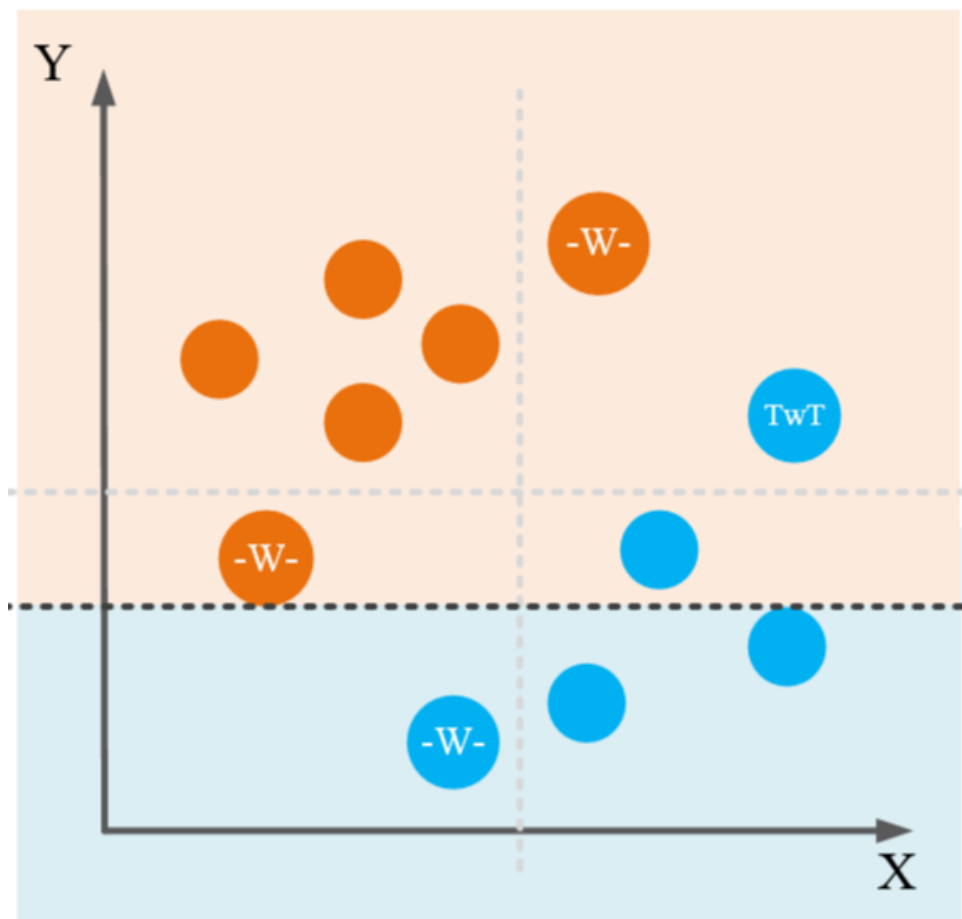


于是再次进行权重调整：




Adaboost

第三次继续划分:




第三次权重调整+第四次继续划分:




Adaboost

最后得到强分类器的决策平面：



Adaboost (MATLAB plot)



Justification

- 令 $Z_m = \sum_{i=1}^N w_i^{(m-1)} \exp[-\alpha_m y_i g_m(x_i)]$, $w_i^{(0)} = \frac{1}{N}$ ($i = 1, 2, \dots, N$),

$$w_i^{(m)} = \frac{w_i^{(m-1)} \exp[-\alpha_m y_i g_m(x_i)]}{Z_m}, G(x) = \sum_{m=1}^M \alpha_m g_m(x).$$

- 对于训练集上的错误率 $R_{\text{tr}}(H_G) = \frac{1}{N} \sum_{i=1}^N |\{i: y_i \neq H_G(x_i)\}|$, 有

$$\begin{aligned} R_{\text{tr}}(H_G) &= \frac{1}{N} \sum_{i=1}^N |\{i: y_i \neq H_G(x_i)\}| = \frac{1}{N} \sum_{i=1}^N \begin{cases} 1 & \text{if } y_i \neq H_G(x_i) \\ 0 & \text{else} \end{cases} \\ &= \frac{1}{N} \sum_{i=1}^N \begin{cases} 1 & \text{if } y_i G(x_i) \leq 0 \\ 0 & \text{else} \end{cases} \\ &\leq \frac{1}{N} \sum_{i=1}^N \exp[-y_i G(x_i)] \\ &= \frac{1}{N} \sum_{i=1}^N \exp[-y_i \sum_{j=1}^m \alpha_j g_j(x_i)] = \prod_{j=1}^m Z_j \end{aligned}$$

Justification

- The training error $R_{tr}(H_G)$ is bounded by $\prod_{j=1}^m Z_j$. In order to minimize the training error $R_{tr}(H_G)$, we can try to minimize Z_m .

$$\begin{aligned}
 \frac{\partial Z_m}{\partial \alpha_m} &= \frac{\partial \sum_{i=1}^N w_i^{(m-1)} \exp[-\alpha_m y_i g_m(x_i)]}{\partial \alpha_m} \\
 &= - \sum_{i=1}^N w_i^{(m-1)} y_i g_m(x_i) \exp[-\alpha_m y_i g_m(x_i)] \\
 &= -\exp(-\alpha_m) \sum_{y_i = g_m(x_i)} w_i^{(m-1)} + \exp(\alpha_m) \sum_{y_i \neq g_m(x_i)} w_i^{(m-1)}.
 \end{aligned}$$

Justification

- From $\frac{\partial Z_m}{\partial \alpha_m} = 0$, we have

$$\exp(-\alpha_m) \sum_{y_i = g_m(x_i)} w_i^{(m-1)} = \exp(\alpha_m) \sum_{y_i \neq g_m(x_i)} w_i^{(m-1)}$$

$$\alpha_m = \frac{1}{2} \ln \frac{\sum_{y_i = g_m(x_i)} w_i^{(m-1)}}{\sum_{y_i \neq g_m(x_i)} w_i^{(m-1)}}$$

- Let $\text{err}_m = \sum_{y_i \neq g_m(x_i)} w_i^{(m-1)}$. Then $\sum_{y_i = g_m(x_i)} w_i^{(m-1)} = 1 - \text{err}_m$

And

$$\alpha_m = \frac{1}{2} \ln \frac{1 - \text{err}_m}{\text{err}_m}$$

Justification

- With $\alpha_m = \frac{1}{2} \ln \frac{1 - \text{err}_m}{\text{err}_m}$ in the following we show that the training error $R_{tr}(H_G)$ decreases exponentially with M .

$$\begin{aligned}
 Z_m &= \sum_{i=1}^N w_i^{(m-1)} \exp[-\alpha_m y_i g_m(x_i)] \\
 &= \exp(-\alpha_m) \sum_{y_i = g_m(x_i)} w_i^{(m-1)} + \exp(\alpha_m) \sum_{y_i \neq g_m(x_i)} w_i^{(m-1)} \\
 &= (1 - \text{err}_{m-1}) \sqrt{\frac{\text{err}_m}{1 - \text{err}_m}} + \text{err}_{m-1} \sqrt{\frac{1 - \text{err}_m}{\text{err}_m}} \\
 &= 2\sqrt{(1 - \text{err}_m)\text{err}_m}.
 \end{aligned}$$

Justification

- Let $\text{err}_m = \frac{1}{2} - \gamma_m$. Since the training error of g on \mathcal{G} is bounded below one half, then $0 < \gamma_m < \frac{1}{2}$. Then we have

$$\begin{aligned} Z_m &= \frac{2\sqrt{(1 - \text{err}_m)\text{err}_m}}{1} \\ &= 2\sqrt{\left(\frac{1}{2} - \gamma_m\right)\left(1 - \frac{1}{2} + \gamma_m\right)} \\ &= \sqrt{1 - 4\gamma_m^2} \\ &\leq \sqrt{\exp(-4\gamma_m^2)} \\ &= \exp(-2\gamma_m^2) \end{aligned}$$

Consider the Taylor expansion of $\exp(-x^2)$ at $x = 0$.

Justification

- Let $\gamma = \min\{\gamma_1, \gamma_2, \dots, \gamma_M\}$. The bound of the training error

$$R_{\text{tr}}(H_G) \leq \prod_{m=1}^M Z_m = \prod_{m=1}^M \exp(-2\gamma_m^2) \leq \prod_{m=1}^M \exp(-2\gamma^2) = \exp(-2\gamma^2 M).$$

- In other words, the training error $R_{\text{tr}}(H_G)$ is bounded by a decaying exponential.

Moreover, since $R_{\text{tr}}(H_G) \in \{0, \frac{1}{N}, \frac{2}{N}, \dots, 1\}$, it follows that after a finite number of steps, when $\exp(-2\gamma^2 M_0) < \frac{1}{N}$, the training error will become 0 and the training data will be perfectly classified!

谢谢各位同学！